GARCH Volatility Analysis - S&P 500

Ian Moore

2025-08-04

GARCH Volatility Modeling

Analysis of S&P 500 Volatility Dynamics

This notebook demonstrates advanced volatility modeling techniques using GARCH models applied to S&P 500 data. We progress from basic univariate GARCH to a multivariate model and rolling forecast. Lastly, we'll compile for risk management applications.

Key Analyses:

1. Univariate GARCH modeling and diagnostics
2. Dynamic Conditional Correlation (DCC-GARCH) for multi-asset portfolios
3. Rolling Window Forecasts
4. Risk management applications (VaR, portfolio optimization)

```
Daily S&P 500 returns were loaded for the period from 2020-01-01 to 2024-12-31.
Additional tickers were loaded for the same period: TLT, GLD, and VXX. These will be
used in the multivariate analysis.
```

1. S&P Summary Statistics and Exploratory Data Analysis

1.1 S&P 500 Summary Statistics and Analysis

```
S&P 500 summary statistics results are the following:

Mean: 0.0628%
Std Dev: 1.3229%
Skewness: -0.5443
Kurtosis: 11.5024
Min: -10.9424%
Max: 9.0603%
```

```
S&P 500 analysis of summary statistics:
```

The S&P 500 daily returns resemble the characteristics of typical financial time series data, with an average daily return of 0.0628%, or approximately 16% annualized. The standard deviation of 1.32% represents moderate daily volatility, while the negative skewness of -0.54 reveals a slight tendency toward larger negative returns than positive returns. The excess kurtosis of 11.50 demonstrates significant fat tails as the baseline kurtosis is 3.0, resulting in an excess of 8.50. This high kurtosis, combined with negative skewness (-0.54), indicates the S&P 500 has frequent small gains, occasional large losses, and reflects extreme events occurring more frequently than normal distribution theory would predict. This is evidenced by the wide range between the observed minimum and maximum from -10.94% to +9.06%. Finally, this confirms the presence of volatility clustering and suggests that traditional risk models assuming normal distributions may underestimate tail risks for this period of the portfolio.

1.2 Exploratory Data Analysis and Visualizations

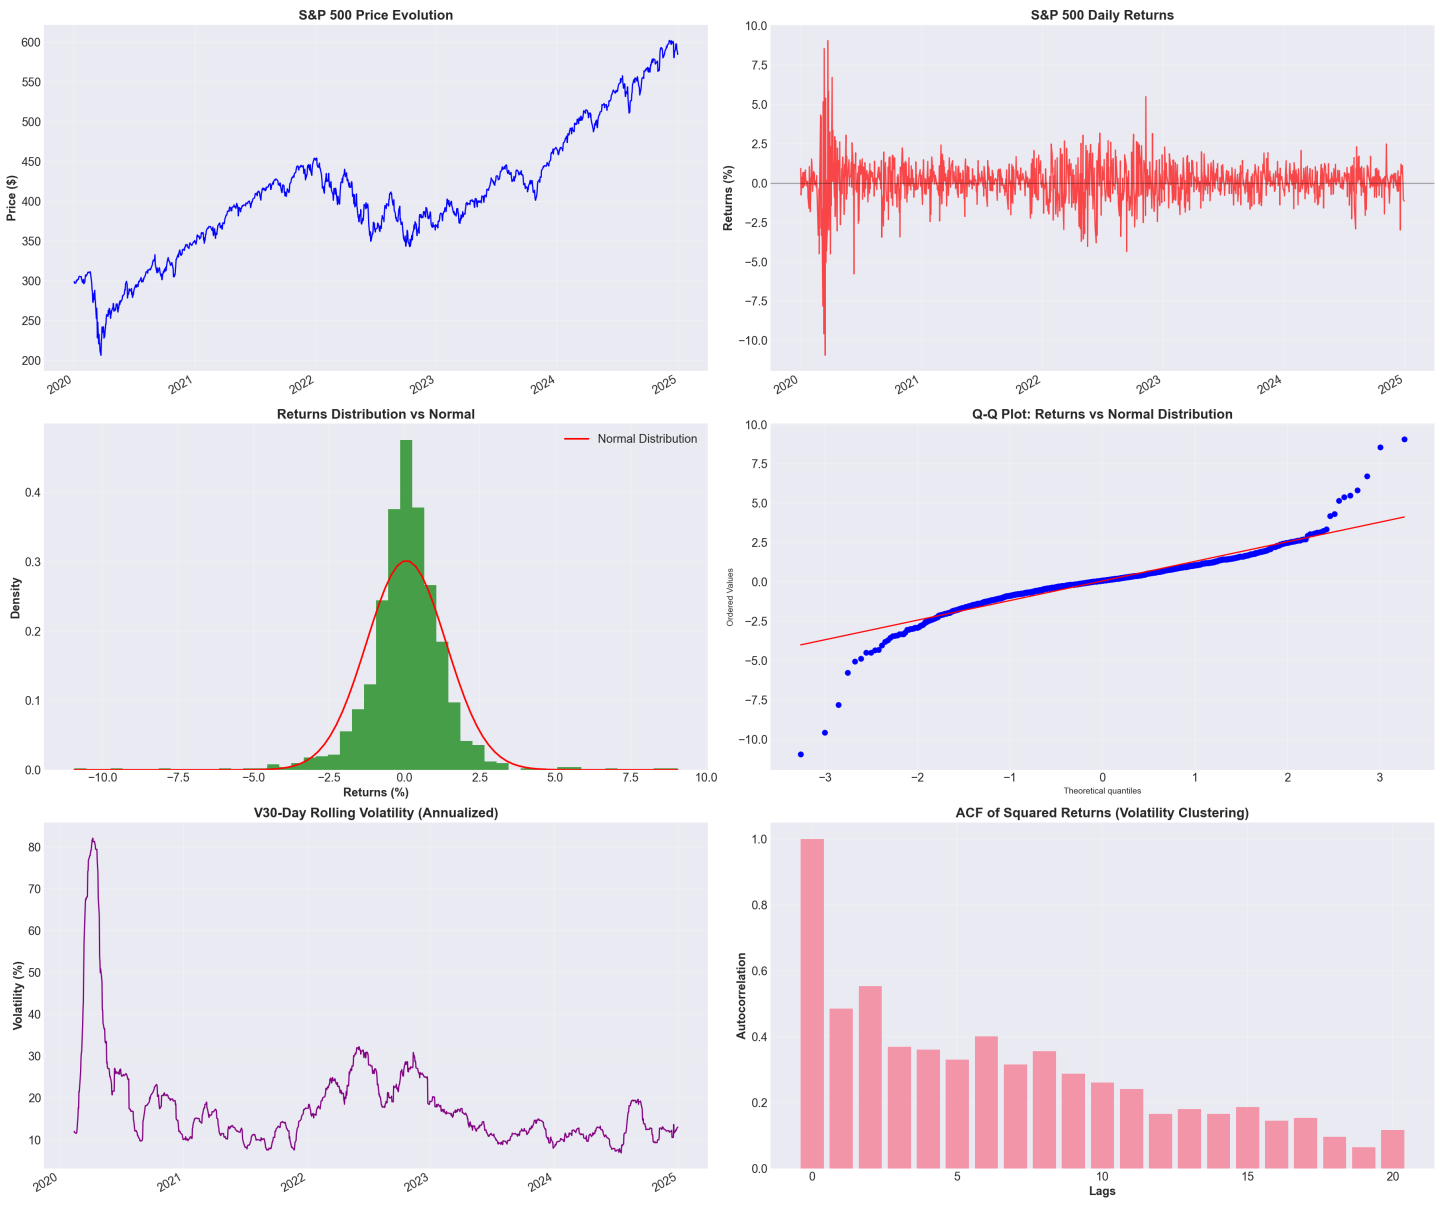

S&P 500 price evolution shows an increase of nearly 200% from $300 to around $600. While demonstrating strong long-term growth, the data reveals two distinct periods of significant volatility clustering: early-to-mid

2020 during the COVID market crash, and throughout most of 2022 amid aggressive interest rate increases. This clustering behavior is clearly visible in the daily returns plot, where periods of high volatility (large price swings) are followed by continued high volatility, and calm periods persist for extended timeframes. The returns distribution and Q-Q plot confirm substantial deviations from normality, particularly in the tails, indicating the presence of extreme market events that occur more frequently than a normal distribution would predict. The 30-day rolling volatility and autocorrelation function (ACF) of squared returns provide deeper insight into this volatility clustering phenomenon. The rolling volatility shows sustained high-volatility periods reaching 80% annualized during COVID and elevated levels throughout 2022, while the ACF demonstrates strong persistence in volatility shocks—when markets become volatile, they tend to remain volatile for weeks rather than quickly reverting to calm conditions. As evidence, the ACF shows volatility remaining elevated even at 20-day lags. This persistent volatility clustering validates the use of GARCH modeling to capture these time-varying risk dynamics that traditional models assuming constant volatility would miss.

1.3 Statistical Tests

```
Statistical Test Results:


Jarque-Bera Test for Normality:
  Statistic: 6930.6696
  P-value: 0.000000
  Result: Reject normality

Ljung-Box Test for ARCH Effects (Volatility Clustering):
  P-value (lag 10): 0.000000
  Result: Significant ARCH effects detected
```

The Jargue-Bera and Ljung-Box tests results appear to be typical and expected for financial data with high volatility clustering. The Jarque-Bera tests rejects normality, which is appropriate given the data's slight skew of -0.54 and fat tails, as indicated in the Q-Q plot. The Ljung-Box test strongly rejects the null hypothesis meaning periods of high volatility tend to be followed by more high volatility, and periods of low volatility tend to be followed by more low volatility. This clustering behavior is visually confirmed in the ACF plot of squared returns, which shows significant autocorrelation persisting for 15-20 days, indicating that volatility shocks have lasting effects rather than quickly reverting to average levels.

2. Univariate GARCH Model

2.1 Model Comparison Test

```
Model Comparison Output (sorted by AIC):
         Model         AIC         BIC  Log-Likelihood  Parameters
  GARCH(1,1)-t 3614.103312 3639.785728    -1802.051656           5
GJR-GARCH(1,1) 3637.198254 3662.880670    -1813.599127           5
    GARCH(1,1) 3658.669768 3679.215701    -1825.334884           4
    GARCH(2,2) 3660.993371 3691.812271    -1824.496686           6
   EGARCH(1,1) 3669.636778 3690.182711    -1830.818389           4

Best model: GARCH(1,1)-t
```

The model comparison results clearly demonstrate that GARCH(1,1) with Student's t-distribution is the optimal choice, achieving the lowest AIC of 3614 compared to 3637 for the next-best model. The substantial improvement from regular GARCH(1,1) (AIC: 3659) to GARCH(1,1)-t (AIC: 3614) confirms the importance of accounting for the fat tails identified in earlier statistical tests. More complex models like GARCH(2,2), GJR-GARCH, and EGARCH failed to provide meaningful improvements despite additional parameters, suggesting that the S&P 500's volatility dynamics follow a relatively simple, symmetric pattern that doesn't require asymmetric or higher-order specifications. This validates the use of GARCH(1,1)-t for volatility forecasting and risk management applications.

```
Model Diagnostics(GARCH(1,1)-t):

                        Constant Mean - GARCH Model Results
==================================================================================
Dep. Variable:                            SPY   R-squared:                   0.000
Mean Model:                     Constant Mean   Adj. R-squared:              0.000
Vol Model:                              GARCH   Log-Likelihood:           -1802.05
Distribution:      Standardized Student's t   AIC:                       3614.10
Method:                    Maximum Likelihood   BIC:                       3639.79
                                                No. Observations:             1257
Date:                        Sun, Aug 03 2025   Df Residuals:                 1256
Time:                                21:09:49   Df Model:                        1
                               Mean Model
========================================================================
                 coef    std err          t      P>|t|    95.0% Conf. Int.
------------------------------------------------------------------------
mu             0.1233  2.314e-02      5.331  9.784e-08 [7.799e-02,  0.169]
                              Volatility Model
==========================================================================
                 coef    std err          t      P>|t|      95.0% Conf. Int.
--------------------------------------------------------------------------
omega          0.0380  1.277e-02      2.978  2.900e-03 [1.300e-02,6.306e-02]
alpha[1]       0.1560  2.872e-02      5.431  5.614e-08   [9.969e-02,  0.212]
beta[1]        0.8247  2.894e-02     28.495 1.356e-178     [  0.768,  0.881]
                              Distribution
========================================================================
                 coef    std err          t      P>|t|  95.0% Conf. Int.
------------------------------------------------------------------------
nu             7.1748      1.408      5.095  3.489e-07 [  4.415,  9.935]
========================================================================

Covariance estimator: robust
```

2.2 Best Model Diagnostics Visualization

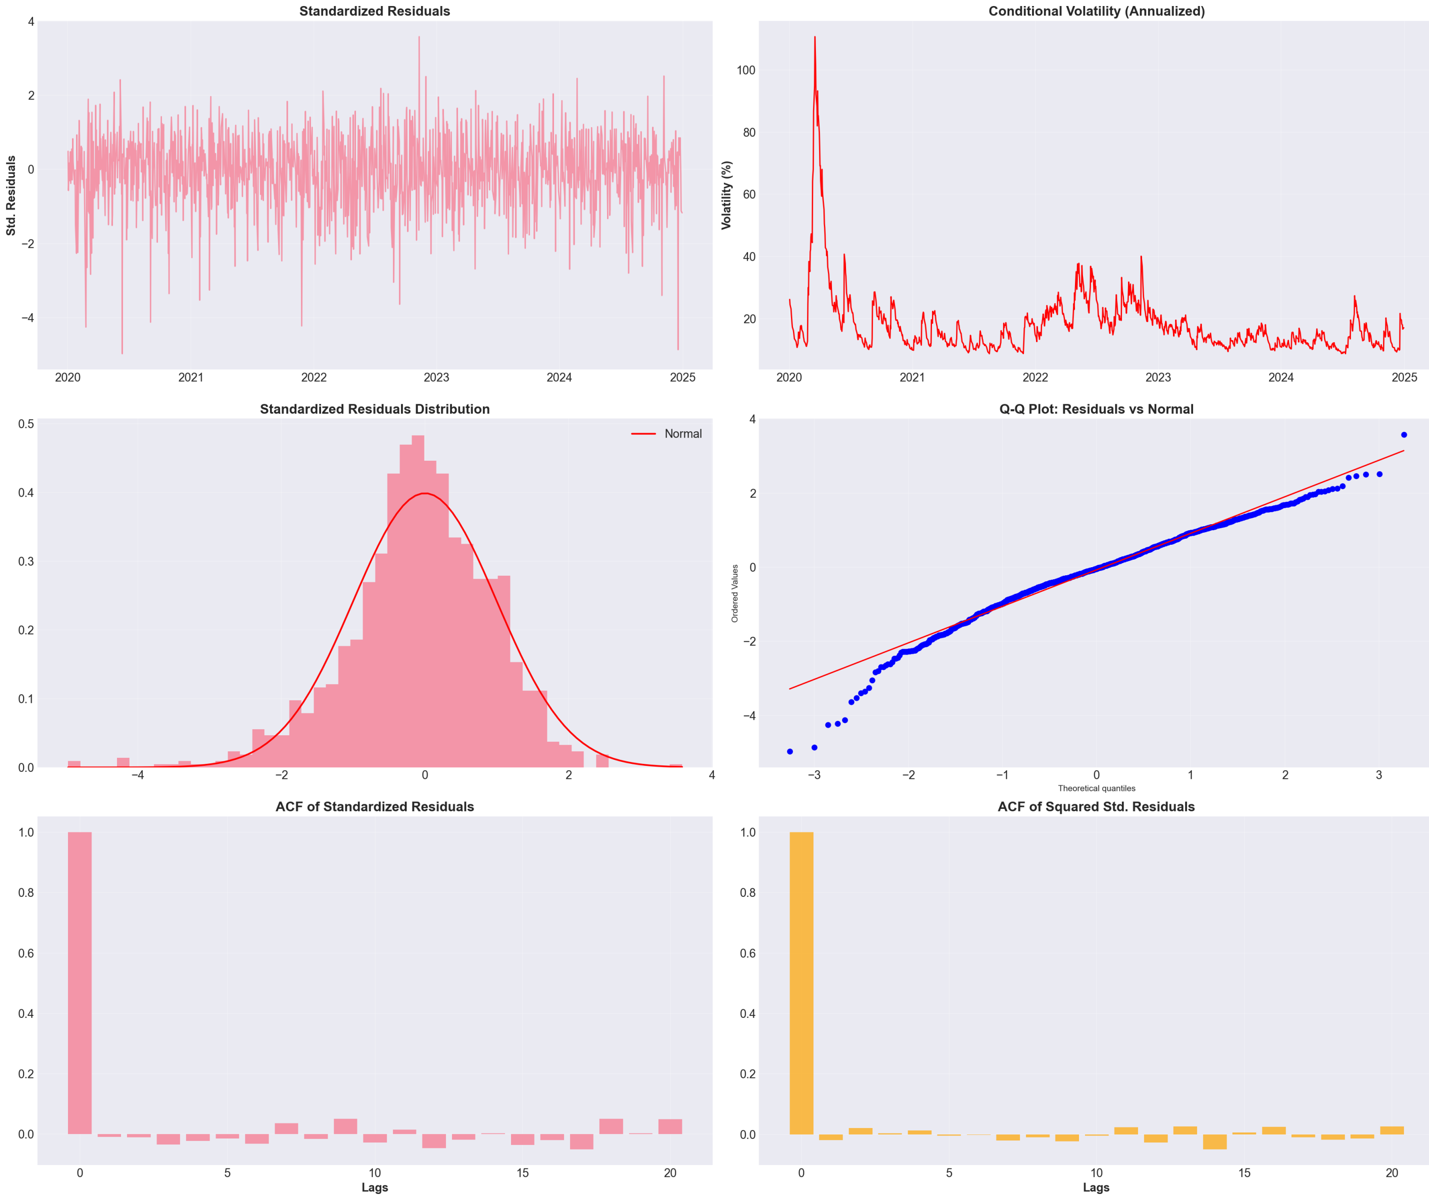

The GARCH(1,1)-t model diagnostics confirm the model is performing across multiple validation criteria. The standardized residuals show no obvious patterns or clustering over time, while the conditional volatility appears to capture the major volatility regimes including the COVID spike, which exceeded 100% annualized volatility, and elevated periods during 2022. The standardized residuals distribution and Q-Q plot demonstrate that while some fat tails remain in the extreme quantiles, the model has substantially improved the distributional fit compared to the raw returns. Most importantly, the ACF plots of both standardized residuals and squared standardized residuals show minimal remaining autocorrelation, indicating the GARCH model has successfully captured the volatility clustering effects that were strongly present in the original data. These diagnostics act as strong validations of the GARCH(1,1)-t models ability to capture the S&P 500's time-varying volatility dynamics.

2.3 Best Model Statistical Tests

```
Residual Diagnostics:

Jarque-Bera (normality): 208.4703 (p=0.000000)
Ljung-Box on squared residuals (ARCH): 0.989380
```

The post-estimation diagnostic tests confirm that the GARCH(1,1)-t model successfully handles the key issues identified in the raw data. While the Jarque-Bera test still rejects normality (statistic: 208.47, p=0.000000), this represents an improvement from the original statistic of 6930.67, and is an indicator the t-distribution has captured most of the fat tails in the return distribution. More importantly, the Ljung-Box test on squared residuals yields a low statistic of 0.989380, suggesting no significant remaining ARCH effects and confirming that the GARCH model has successfully eliminated the volatility clustering that was strongly present in the original data. These results validate that the model effectively captures the time-varying volatility dynamics while maintaining realistic distributional assumptions for financial returns.

3. Multivariate GARCH Model

3.1 Univariate Model Evaluation for Comparison and Correlation

`As indicated above, additional tickers were loaded for evaluation: TLT, GLD, and VXX.`

Need commentary here

3.2 Crisis Analysis Correlations

```
Correlation During High VIX Periods (VIX > 2.4)):

  SPY_TLT: Crisis=-0.080, Normal=-0.034, Difference=-0.046
  SPY_GLD: Crisis=0.140, Normal=0.155, Difference=-0.015
  TLT_GLD: Crisis=0.303, Normal=0.304, Difference=-0.001
```

Add commentary here

3.3 Conditional Volatility and Dynamic Coorelation Evaluation

```
Average Portfolio Correlations(SPY, TLT, GLD):

  SPY-TLT: -0.043
  SPY-GLD: 0.152
  TLT-GLD: 0.304
```

Short commentary goes here

```
Portfolio Optimization and GARCH Comparison:
```

```
Portfolio Optimization Using Dynamic Correlations:
Average Portfolio Volatility (DCC): 11.00%
Average Portfolio Volatility (Historical): 10.66%
```

Comment goes here

```
Portfolio Volatility Statistics and Visualization:
```

```
Number of observations: 1257
Date range: 2020-01-02 to 2024-12-30
DCC Portfolio Volatility:
  Mean: 11.00%
  Std:  3.86%
  Min:  7.26%
  Max:  42.56%
```

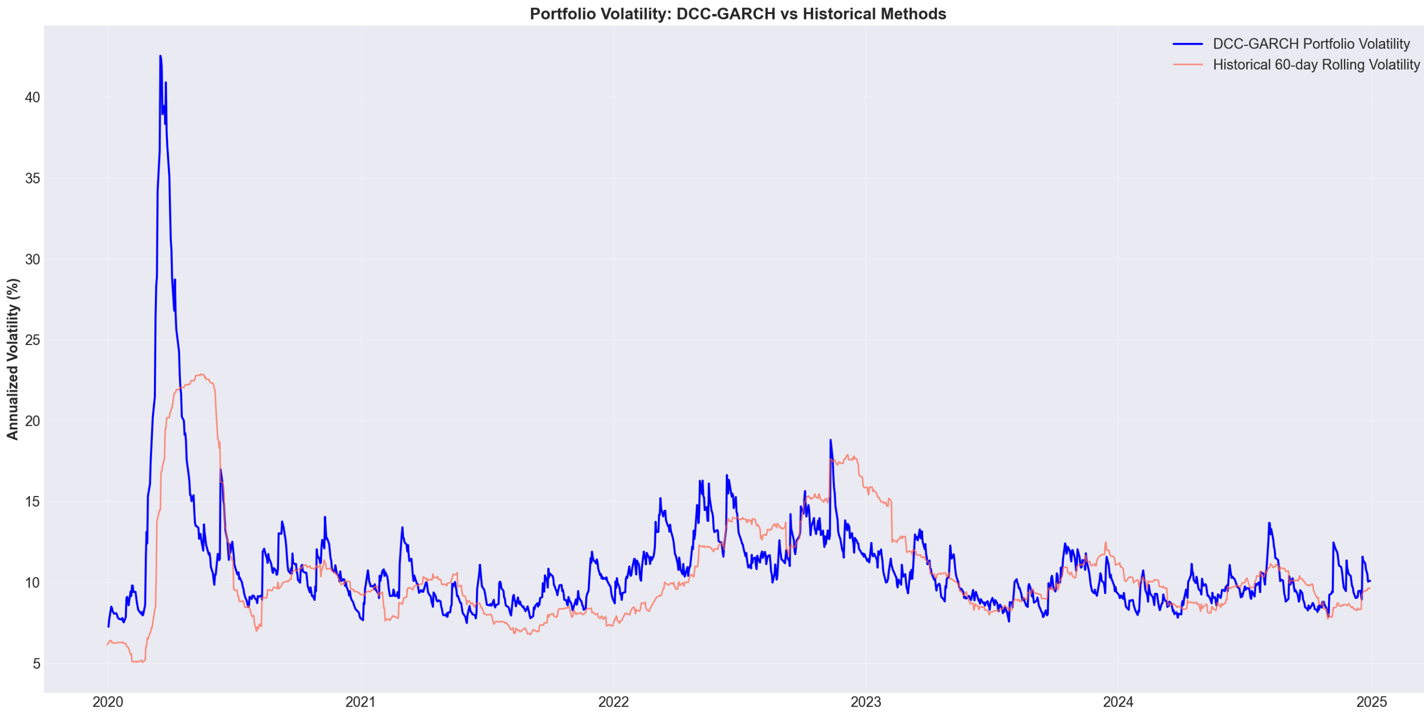

Conclusionary commentary here on multivaraite analysis

4. Rolling Window Forecast

```
Simple Rolling GARCH Forecast
Window size: 252 days
Forecast horizon: 30 days
Forecasting every 10 days



  Generated 98 forecasts


Forecast Accuracy Metrics:
  Mean Absolute Error (MAE): 3.920%
  Root Mean Square Error (RMSE): 4.923%
  Mean Absolute Percentage Error (MAPE): 26.46%
  Correlation: 0.6864
  Direction Accuracy: 100.00%
  Number of Observations: 98
```

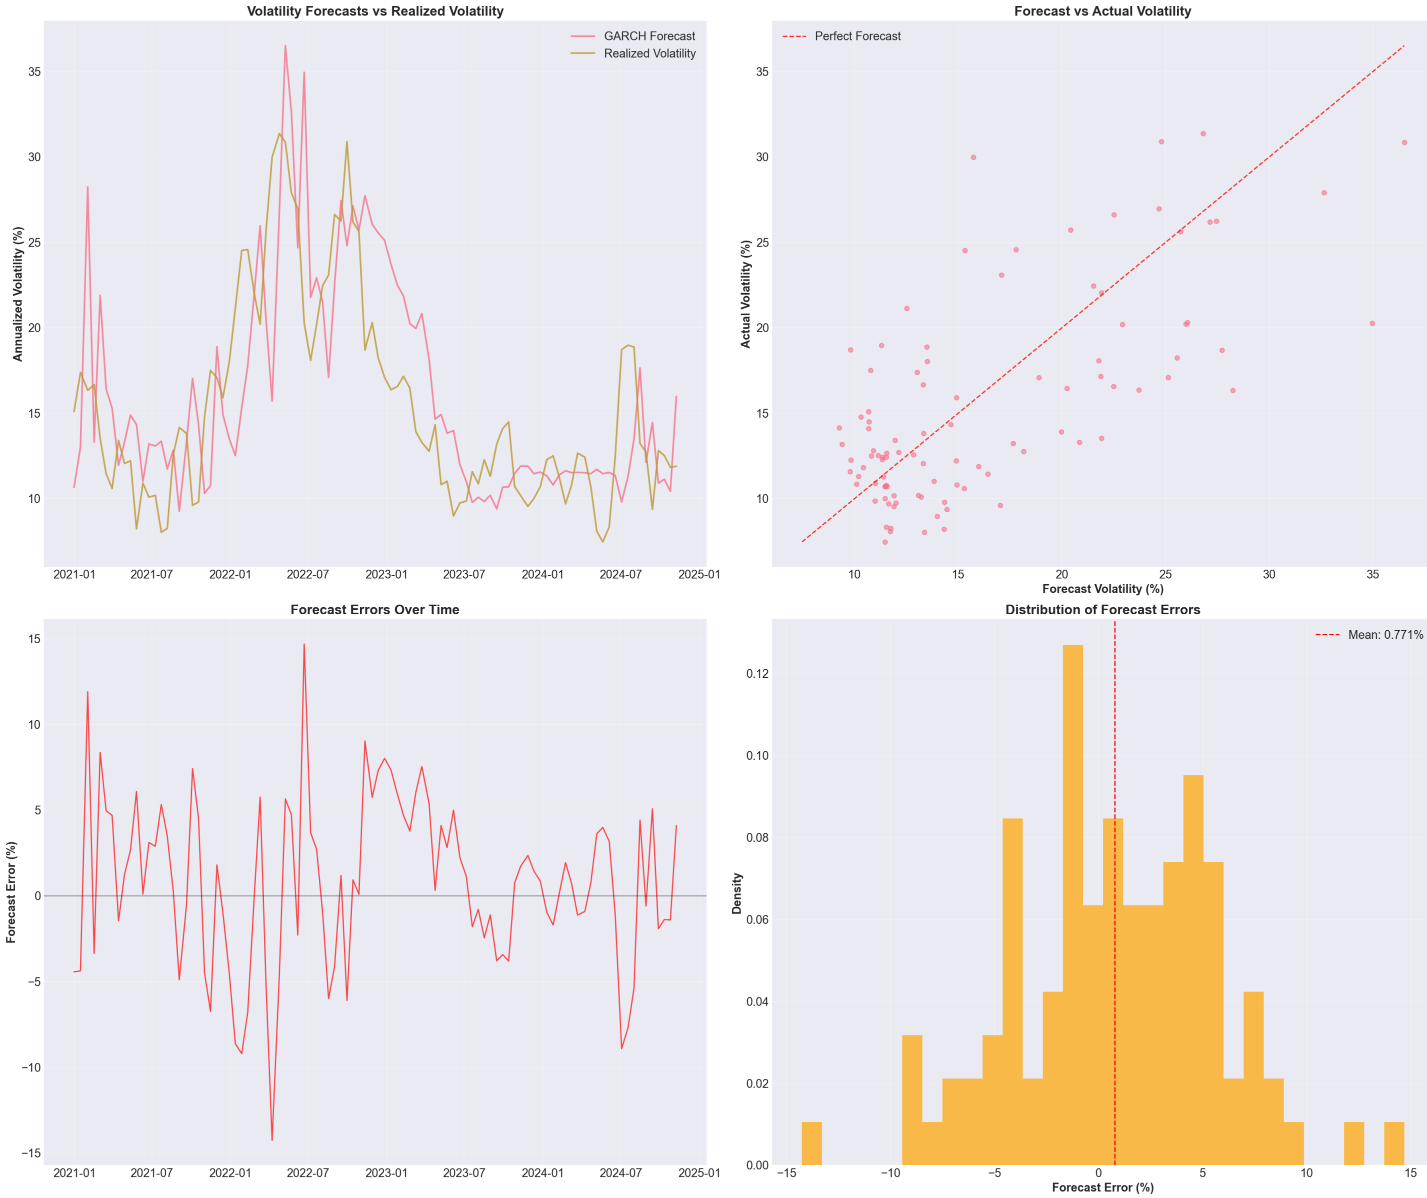

5. RISK MANAGEMENT APPLICATIONS

```
Calculating Dynamic VaR (1-day horizon)
Confidence levels: [95, 99]%
Distribution: t
--------------------------------------------------
Volatility forecast: 1.117%
95% VaR: -2.171%
99% VaR: -3.511%
--------------------------------------------------

Portfolio VaR (Portfolio Value: $1,000,000)
--------------------------------------------------
95% VaR: -2.171% = $21,708
99% VaR: -3.511% = $35,109
--------------------------------------------------
```

```
VaR Summary Table:
============================================================
Confidence Level VaR (%) Volatility Forecast (%) VaR ($)
             95% -2.171%                  1.117% $21,708
             99% -3.511%                  1.117% $35,109
============================================================

Interpretation:
• 5% chance of daily loss exceeding 2.17% ($21,708)
• 1% chance of daily loss exceeding 3.51% ($35,109)

STRESS TESTING SCENARIOS
============================================================
Portfolio Value: $1,000,000
============================================================

STRESS TEST RESULTS
--------------------------------------------------------------------------------

BASE CASE:
  Description: Normal market conditions (1-day GARCH forecast)
  Annual Volatility: 17.7%
  95% Daily VaR: -2.17% ($21,708)
  99% Daily VaR: -3.51% ($35,109)
  Expected Annual Loss: 7.1% ($70,937)

MODERATE STRESS:
  Description: 1.5x current market volatility
  Annual Volatility: 19.6%
  95% Daily VaR: -38.02% ($380,175)
  99% Daily VaR: -61.49% ($614,850)
  Expected Annual Loss: 7.8% ($78,258)

CRISIS 2008:
  Description: 2008-style financial crisis (45% volatility)
  Annual Volatility: 45.0%
  95% Daily VaR: -87.44% ($874,431)
  99% Daily VaR: -141.42% ($1,414,201)
  Expected Annual Loss: 27.0% ($270,000)

EXTREME COVID:
  Description: COVID-style market shock (60% volatility)
  Annual Volatility: 60.0%
  95% Daily VaR: -116.59% ($1,165,908)
  99% Daily VaR: -188.56% ($1,885,601)
  Expected Annual Loss: 48.0% ($480,000)

================================================================================

Creating GARCH Risk Management Dashboard...
<arch.univariate.base.ARCHModelForecast object at 0x000002505E5F6690>
```

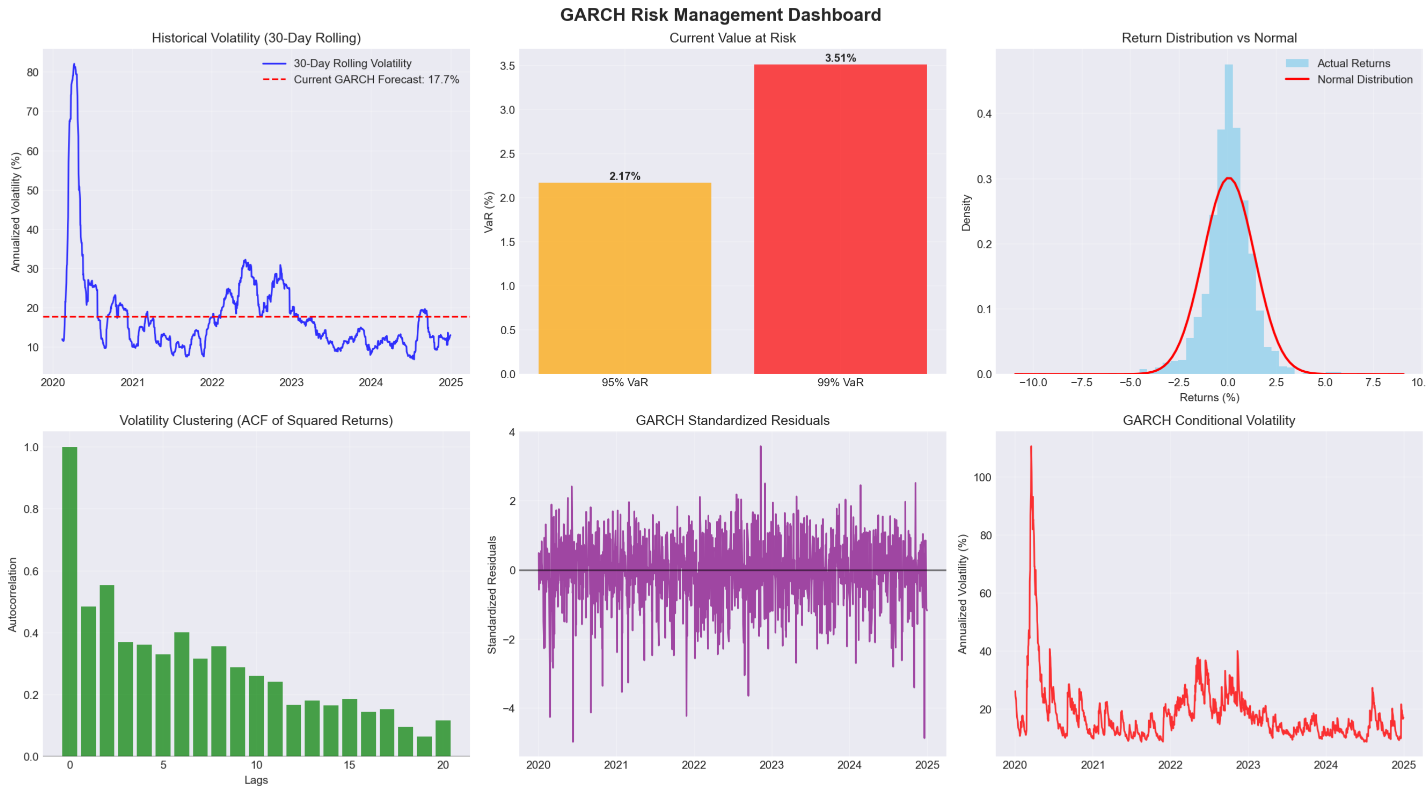


```
============================================================
RISK DASHBOARD SUMMARY
============================================================
Current Volatility Forecast: 17.73%
30-Day Realized Volatility: 13.04%
1-Year Realized Volatility: 12.57%

95% VaR (1-day): 2.17%
99% VaR (1-day): 3.51%

Risk Assessment:
  MODERATE RISK: Current volatility elevated
============================================================



==================================================
PORTFOLIO RISK SUMMARY
==================================================
Individual Asset Volatilities:
  SPY: 21.00% (Weight: 25.0%)
  TLT: 17.97% (Weight: 25.0%)
  GLD: 15.53% (Weight: 25.0%)
  VXX: 75.92% (Weight: 25.0%)

Portfolio Volatility: 17.27%

Correlation Matrix:
       SPY    TLT    GLD    VXX
```

```
SPY  1.000 -0.150  0.158 -0.710
TLT -0.150  1.000  0.267  0.129
GLD  0.158  0.267  1.000 -0.095
VXX -0.710  0.129 -0.095  1.000

Diversification Analysis:
  Weighted Average Volatility: 32.60%
  Portfolio Volatility: 17.27%
  Diversification Benefit: 15.33%
==================================================
```